

```
# prompt: mount drive
```

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
import numpy as np
```

```
# Load the data from the .npz file
data = np.load('/content/rotated_mnist (1).npz')
```

```
# Access the data arrays
rotated_x_train = data['x_train']
rotated_y_train = data['y_train']
rotated_x_test = data['x_test']
rotated_y_test = data['y_test']
```

```
# Print the shape of the data arrays to show the data clust
print("Shape of rotated_x_train:", rotated_x_train.shape)
print("Shape of rotated_y_train:", rotated_y_train.shape)
print("Shape of rotated_x_test:", rotated_x_test.shape)
print("Shape of rotated_y_test:", rotated_y_test.shape)
```

Shape of rotated\_x\_train: (152400, 28, 28)  
 Shape of rotated\_y\_train: (152400,)  
 Shape of rotated\_x\_test: (26004, 28, 28)  
 Shape of rotated\_y\_test: (26004,)

```
!pip install tensorflow
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
```

Requirement already satisfied: tensorflow in /usr/local/lib/python3.11/dist-packages (2.18.0)  
 Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.4.0)  
 Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.6.3)  
 Requirement already satisfied: flatbuffers>=24.3.25 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (25.2.10)  
 Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>=0.2.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.6.0)  
 Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.2.0)  
 Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (18.1.1)  
 Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.4.0)  
 Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from tensorflow) (24.2)  
 Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!4.21.3,!4.21.4,!4.21.5,<6.0.0dev,>=3.20.3 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.21.0)  
 Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (2.32.3)  
 Requirement already satisfied: setuptools in /usr/local/lib/python3.11/dist-packages (from tensorflow) (75.2.0)  
 Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.17.0)  
 Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.0.0)  
 Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (4.13.0)  
 Requirement already satisfied: wrapt>=1.11.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.17.2)  
 Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.71.0)  
 Requirement already satisfied: tensorboard<2.19,>=2.18 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (2.18.0)  
 Requirement already satisfied: keras>=3.5.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.8.0)  
 Requirement already satisfied: numpy<2.1.0,>=1.26.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (2.0.2)  
 Requirement already satisfied: h5py>=3.11.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.13.0)  
 Requirement already satisfied: ml-dtypes<0.5.0,>=0.4.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.4.1)  
 Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.37.0)  
 Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.11/dist-packages (from astunparse>=1.6.0->tensorflow) (0.44.0)  
 Requirement already satisfied: rich in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (13.9.4)  
 Requirement already satisfied: nameex in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (0.0.8)  
 Requirement already satisfied: optree in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (0.14.1)  
 Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.4.0)  
 Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.10.1)  
 Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (2.3.0)  
 Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (2025.1.31)  
 Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2.19,>=2.18->tensorflow) (3.7.0)  
 Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2.19,>=2.18->tensorflow) (0.6.0)  
 Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2.19,>=2.18->tensorflow) (3.1.0)  
 Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.11/dist-packages (from werkzeug>=1.0.1->tensorboard<2.19,>=2.18->tensorflow) (3.0.2)  
 Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras>=3.5.0->tensorflow) (3.0.0)  
 Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras>=3.5.0->tensorflow) (2.19.0)  
 Requirement already satisfied: mdurl~0.1 in /usr/local/lib/python3.11/dist-packages (from markdown-it-py>=2.2.0->rich->keras>=3.5.0->tensorflow) (0.1.2)

## > Loading pre-trained vae model

[ ] ↳ 2 cells hidden

## › Supervised ORACLE MODEL

[ ] ↳ 9 cells hidden

## ✓ Unsupervised ORACLE MODEL

```
_, _, z_train = encoder.predict(rotated_x_train)
_, _, z_test = encoder.predict(rotated_x_test)

latent_dim = z_train.shape[1]

class SymmetryGenerator(keras.Model):
    def __init__(self, latent_dim):
        super(SymmetryGenerator, self).__init__()
        self.net = keras.Sequential([
            layers.Dense(64, activation='relu'),
            layers.Dense(64, activation='relu'),
            layers.Dense(latent_dim)
        ])

    def call(self, inputs):
        return self.net(inputs)

def invariance_loss(z, generator, classifier, epsilon=1e-3):
    z_prime = z + epsilon * generator(z)
    return tf.reduce_mean(tf.square(classifier(z) - classifier(z_prime)))

def normalization_loss(z, generator):
    norms = tf.norm(generator(z), axis=1)
    mean_norm = tf.reduce_mean(norms)
    return (
        tf.reduce_mean(tf.square(norms - 1.0)) +
        tf.reduce_mean(tf.square(norms - mean_norm))
    )

def orthogonality_loss(generators, z):
    loss = 0
    for i in range(len(generators)):
        for j in range(i+1, len(generators)):
            dot_prods = tf.reduce_sum(generators[i](z) * generators[j](z), axis=1)
            loss += tf.reduce_mean(tf.square(dot_prods))
    return loss

def build_classifier(latent_dim):
    classifier = keras.Sequential([
        layers.Input(shape=(latent_dim,)),
        layers.Dense(128, activation='relu'),
        layers.Dense(128, activation='relu'),
        layers.Dense(32, activation='relu'),
        layers.Dense(1, activation='sigmoid') #Output layer is 1, as we are predicting the difference between z and z'
    ])
    classifier.compile(
        optimizer=keras.optimizers.Adam(learning_rate=1e-3),
        loss='mse' #Using mean squared error as we are training to predict the difference between z and z'
    )
    return classifier

↳ 4763/4763 ————— 8s 1ms/step
813/813 ————— 2s 2ms/step

def train_symmetry_generator(z, num_generators=1):
    generators = [SymmetryGenerator(latent_dim) for _ in range(num_generators)]
    optimizer = keras.optimizers.Adam(learning_rate=3e-4)
    classifier = build_classifier(latent_dim)

    for epoch in range(50):
        with tf.GradientTape() as tape:
            inv_loss = sum(invariance_loss(z, gen, classifier) for gen in generators)
            norm_loss = sum(normalization_loss(z, gen) for gen in generators)
            ortho_loss = orthogonality_loss(generators, z) if num_generators > 1 else 0

            total_loss = inv_loss + norm_loss + ortho_loss
```

```

grads = tape.gradient(total_loss, [g.trainable_variables for g in generators])
for gen, grad in zip(generators, grads):
    optimizer.apply_gradients(zip(grad, gen.trainable_variables))

#Train classifier to predict the difference between z and z'
z_prime = z + 1e-3 * generators[0](z) #Using the first generator for training.
classifier.train_on_batch(z, z-z_prime)

```

```
return generators, classifier
```

```
generators, classifier = train_symmetry_generator(z_train, num_generators=1)
```

```
def generate_symmetric_samples(z, generator, epsilon=0.1):
    return z + epsilon * generator(z)
```

```
augmented_z = generate_symmetric_samples(z_train, generators[0])
augmented_images = decoder.predict(augmented_z)
```

↩ 4763/4763 ————— 7s 1ms/step

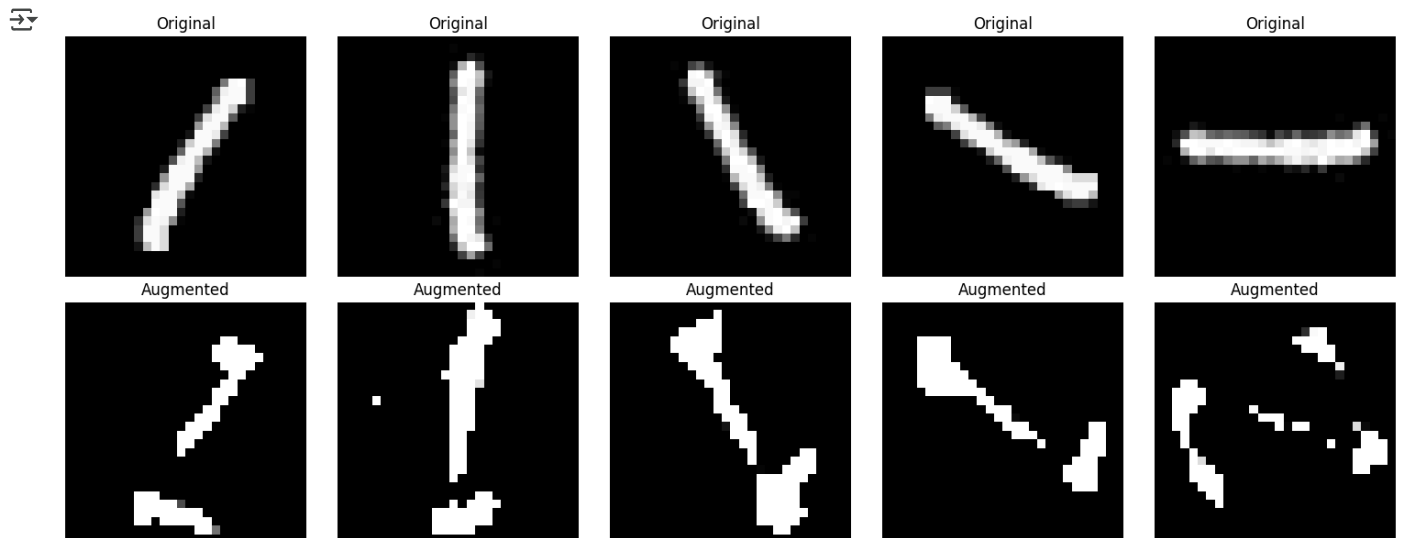
```
import matplotlib.pyplot as plt
```

```
def visualize_symmetry_effects(original_images, augmented_images, num_samples=5):
    plt.figure(figsize=(15, 6))
    for i in range(num_samples):
        plt.subplot(2, num_samples, i+1)
        plt.imshow(original_images[i].squeeze(), cmap='gray')
        plt.title("Original")
        plt.axis('off')

        plt.subplot(2, num_samples, num_samples+i+1)
        plt.imshow(augmented_images[i].squeeze(), cmap='gray')
        plt.title("Augmented")
        plt.axis('off')

    plt.tight_layout()
    plt.show()
```

```
visualize_symmetry_effects(rotated_x_train[:5], augmented_images[:5])
```



```

# Apply symmetry transformation to latent space
epsilon = 2 # Can be adjusted till 2 distortion
z_transformed = z_train + epsilon * generators[0](z_train)

```

```

# Decoding the transformed latent representations
decoded_images = decoder.predict(z_transformed)

```

↩ 4763/4763 ————— 6s 1ms/step

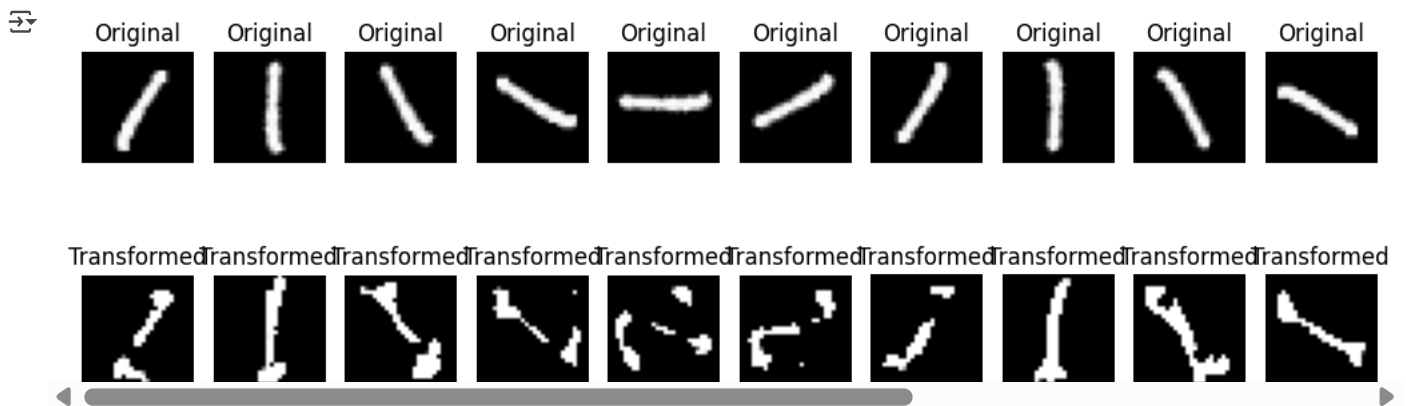
```
import matplotlib.pyplot as plt

num_images = 10

plt.figure(figsize=(10, 4))
for i in range(num_images):
    # Original image
    plt.subplot(2, num_images, i + 1)
    plt.imshow(rotated_x_train[i].squeeze(), cmap='gray')
    plt.axis('off')
    plt.title("Original")

    # Transformed image
    plt.subplot(2, num_images, num_images + i + 1)
    plt.imshow(decoded_images[i].squeeze(), cmap='gray')
    plt.axis('off')
    plt.title("Transformed")

plt.tight_layout()
plt.show()
```



## ✓ Manually finding the indice of 2

```
import matplotlib.pyplot as plt

def generate_symmetric_samples(z, generator, epsilon=0.1):
    return z + epsilon * generator(z)


image_indices = [ 91981,91982,91983]

for i in image_indices:
    augmented_z = generate_symmetric_samples(z_train[i:i+1], generators[0], epsilon=2) # symmetry transformations with epsilon = 2
    augmented_image = decoder.predict(augmented_z)

    plt.figure()
    plt.imshow(augmented_image[0].squeeze(), cmap='gray')
    plt.title(f"Augmented Image of '2' - Example {i+1}")
    plt.axis('off')
    plt.show()
```

1/1 — 0s 35ms/step

Augmented Image of '2' - Example 91982



1/1 — 0s 34ms/step

Augmented Image of '2' - Example 91983



1/1 — 0s 34ms/step

Augmented Image of '2' - Example 91984



```
import matplotlib.pyplot as plt
```

```
image_indices = [ 91981,91982,91983]
```

```
for i in image_indices:  
    # Original image  
    plt.figure()  
    plt.imshow(rotated_x_train[i].squeeze(), cmap='gray')  
    plt.title(f"Original Image - Example {i+1}")  
    plt.axis('off')  
    plt.show()
```

```
augmented_z = generate_symmetric_samples(z_train[i:i+1], generators[0], epsilon=2)
augmented_image = decoder.predict(augmented_z)

plt.figure()
plt.imshow(augmented_image[0].squeeze(), cmap='gray')
plt.title(f"Augmented Image - Example {i+1}")
plt.axis('off')
plt.show()
```




Original Image - Example 91982




1/1 — 0s 37ms/step

Augmented Image - Example 91982



Original Image - Example 91983



Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.